MaAI

A **Real-time** and **Light-weight** Software for Generation of **Non-Linguistic** Behavior Generations for Conversational AIs

(Real-time Implementation of Voice Activity Projection)

README: English | Japanese (日本語)

MaAI is a state-of-the-art and light-weight software that can generate (predict) non-linguistic behaviors in real time and continuously. It supports essential interaction elements such as **turn-taking**, **backchanneling**, and **nodding**. Currently available for English GB, Chinese CN, and Japanese JP languages, MaAI will continue to expand its language coverage and non-linguistic behavior repertoire in the future. Designed specifically for conversational AI, including spoken dialogue systems and interactive robots, MaAI handles audio input effectively in either two-channels (user-system) or single-channel (user-only) settings. Thanks to its lightweight design, MaAI operates efficiently, even exclusively on CPU hardware.

The currently supported models are mainly based on the Voice Activity Projection (VAP) model and its extensions. Details about the VAP model can be found in the following repository: https://github.com/ErikEkstedt/VoiceActivityProjection

Update

- We launched the MaAI project and repository here! (July 14th, 2024)

Getting Started

To quickly get started with MaAI, you can install it using pip:

```
pip install maai
```

> **Note:** By default, the CPU version of PyTorch will be installed. If you wish to run MaAI on a GPU, please install the GPU version of PyTorch that matches your CUDA environment before proceeding.

You can run it as follows. The appropriate model for the task (mode) and parameters will be downloaded automatically.
Below is an example where two wav files (user and system) are input to the turn-taking model (VAP).

```
from maai import Maai, MaaiInput

wav1 = MaaiInput.Wav(wav_file_path="path_to_your_user_wav_file")
wav2 = MaaiInput.Wav(wav_file_path="path_to_your_system_wav_file")

maai = Maai(mode="vap", frame_rate=10, context_len_sec=5, audio_ch1=wav1,
audio_ch2=wav2, device="cpu")

maai.start_process()
while True:
    result = maai.get_result()
```

Models

We support the following models (behavior, language, audio setting, etc.), and more models will be added in the future.

Turn-Taking

The turn-taking model uses the original VAP as is and predicts which participant will speak in the next moment.

- VAP Model
  - Japanese
  - English
  - Chinese
- Noise-Robusst VAP Model (**Recommended**)
  - Japanese
- Single-Channel VAP Model
  - In Preparation ...

Backchannel

Backchannels are short listener responses such as `yeah` and `oh`, that are also related to turn-taking.

- VAP-based Backchannel Prediction Model
  - Japanese - Timing Only
  - Japanese - Timing for Two types
- Noise-Robusst VAP-BC
  - In Preparation ...
- Single-Channel VAP-BC
  - In Preparation ...

Nodding

Nodding refers to the up-and-down movement of the head and is closely related to backchanneling. Unlike backchannels that involve vocal responses, nodding allows the listener to express their reaction non-verbally.

- VAP-based Nodding Prediction Model
  - Japanese

Input / Output

For input to the MaAI model, you can directly call the `process` method of a `Maai` class instance.
The `MaaiInput` class also provides flexible input options, supporting audio from WAV files, microphone input, and TCP communication.

- WAV file input: `Wav` class
- Microphone input: `Mic` class
- TCP communication: `TCPReceiver` / `TCPTransmitter` classes

By using these classes, you can easily adapt the audio input method to your specific use case.

For output, the `MaaiOutput` class is currently under development.
At present, you can retrieve the processing results using the `get_result` method of the `Maai` class instance.

For more details, please refer to the following README files:

- Input Readme
- Output Readme

Example Implementation

You can find example implementations of MaAI models in the test_scripts directory of this repository.

- Turn-Taking (VAP)
    - With 2 wav file inputs 🎧
    - With 2 mic inputs 🎤
    - With 2 mic inputs via TCP networks 🌐
    - With 1 wav file adn 1 mic inputs 🎧🎤
- Backchannel
    - With 1 wav file adn 1 mic inputs 🎧🎤
- Nodding
    - With 1 wav file adn 1 mic inputs 🎧🎤

Publication

Please cite the following paper, if you made any publications made with this repository. 🙏

Koji Inoue, Bing'er Jiang, Erik Ekstedt, Tatsuya Kawahara, Gabriel Skantze
Real-time and Continuous Turn-taking Prediction Using Voice Activity Projection
International Workshop on Spoken Dialogue Systems Technology (IWSDS), 2024
https://arxiv.org/abs/2401.04868

```
@inproceedings{inoue2024iwsds,
    author = {Koji Inoue and Bing'er Jiang and Erik Ekstedt and Tatsuya Kawahara
and Gabriel Skantze},
    title = {Real-time and Continuous Turn-taking Prediction Using Voice Activity
Projection},
    booktitle = {International Workshop on Spoken Dialogue Systems Technology
(IWSDS)},
    year = {2024},
    url = {https://arxiv.org/abs/2401.04868},
}
```

If you use the multi-lingual VAP model, please also cite the following paper.

Koji Inoue, Bing'er Jiang, Erik Ekstedt, Tatsuya Kawahara, Gabriel Skantze
Multilingual Turn-taking Prediction Using Voice Activity Projection
Joint International Conference on Computational Linguistics, Language Resources and Evaluation (LREC-COLING), pages 11873-11883, 2024
https://aclanthology.org/2024.lrec-main.1036/

```
@inproceedings{inoue2024lreccoling,
    author = {Koji Inoue and Bing'er Jiang and Erik Ekstedt and Tatsuya Kawahara
and Gabriel Skantze},
    title = {Multilingual Turn-taking Prediction Using Voice Activity Projection},
```

```
    booktitle = {Proceedings of the Joint International Conference on
Computational Linguistics and Language Resources and Evaluation (LREC-COLING)},
    pages = {11873--11883},
    year = {2024},
    url = {https://aclanthology.org/2024.lrec-main.1036/},
}
```

If you also use the noise-robusst VAP model, please also cite the following paper.

Koji Inoue, Yuki Okafuji, Jun Baba, Yoshiki Ohira, Katsuya Hyodo, Tatsuya Kawahara
A Noise-Robust Turn-Taking System for Real-World Dialogue Robots: A Field Experiment
https://www.arxiv.org/abs/2503.06241

```
@misc{inoue2025noisevap,
    author = {Koji Inoue and Yuki Okafuji and Jun Baba and Yoshiki Ohira and
Katsuya Hyodo and Tatsuya Kawahara},
    title = {A Noise-Robust Turn-Taking System for Real-World Dialogue Robots: A
Field Experiment},
    year = {2025},
    note = {arXiv:2503.06241},
    url = {https://www.arxiv.org/abs/2503.06241},
}
```

If you also use the backchannel VAP model, please also cite the following paper.

Koji Inoue, Divesh Lala, Gabriel Skantze, Tatsuya Kawaharaa
Yeah, Un, Oh: Continuous and Real-time Backchannel Prediction with Fine-tuning of Voice Activity Projection
https://aclanthology.org/2025.naacl-long.367/

```
@inproceedings{inoue2025vapbc,
    author = {Koji Inoue and Divesh Lala and Gabriel Skantze and Tatsuya
Kawahara},
    title = {Yeah, Un, Oh: Continuous and Real-time Backchannel Prediction with
Fine-tuning of Voice Activity Projection},
    booktitle = {Proceedings of the Conference of the Nations of the Americas
Chapter of the Association for Computational Linguistics: Human Language
Technologies (NAACL)},
    pages = {7171--7181},
    year = {2025},
    url = {https://aclanthology.org/2025.naacl-long.367/},
}
```

License

The source code in this repository is licensed under the MIT license. The trained models, found in the asset directory, are used for only academic purposes.

A pre-trained CPC model, located at `asset/cpc/60k_epoch4-d0f474de.pt`, is from the original CPC project and please follow its specific license. Refer to the original repository at https://github.com/facebookresearch/CPC_audio for more details.